

OCaml as fast as C

Sylvain Le Gall, OCamlCore SARL,
OCaml Meeting 2009 – Grenoble, France

Plan

- Project workflow
- Optimize
- Tips and tricks

Foreword

This presentation is based on the project STSort which has been funded Talend.

- GNU sort like utility (POSIX/Windows)
- Performances
- Features
- 6 months
- 1 man
- 25 kLoC
- 10% of C

Goal

- This presentation will try to show how:
 - To integrate testing/benchmarking
 - To manage a project « designed for speed »
 - To convince your boss that OCaml is the right choice to do that

**However this presentation is only self tested,
most of the goal above are quite difficult to
achieve (in particular the last point).**

OCaml choice (criticism)

- OCaml is not a widely known
- Your boss prefer Java/C#/C++
- Documentation and examples are hard to find
- Garbage collection is slow

**For all this reasons, your boss or even you,
will avoid using OCaml for speed**

This is an error!

OCaml choice

- Most of the time, code written in Java/COBOL/C++ won't never been read another time, if it works.
- OCaml allow to catch enough error to get a working program that nobody will read again
- OCaml and functional programming community is growing
- OCaml GC is efficient and can achieve a good speed


**For all these reasons you can choose
OCaml for speed**

OCaml choice

- Multi-paradigm language
- Strict typing
- Compile-time checks
- Fast prototyping
- Static compilation
- Spend less time finding NULL pointer error
- Develop application faster

If OCaml is your most fluent programming language, just choose it

Programming



- OCaml allow fast prototyping:
 - Establish a very short circuit with customer
 - Find the root of what they need
 - Show them something that works quickly

When prototyping, « quick and dirty » is the rule of thumb

Programming (after prototype)

- Stop and think:
 - What is really useful in my prototype?
 - What can I reuse in it?
 - Are there things I can replace with Open Source libraries?
 - Draw the general architecture of your software
- Evaluate your need in term of time and workforce
- Begin to gather information on what has failed in your prototype!

Programming

- Try to deliver weekly/monthly, even if this is not finished, it gives you early feedback and show your progress (including to you)
- Use a good build tool:
 - OcamlMakefile: portable, small project
 - Omake: large project
 - Ocamlbuild: ≥ 3.10
- Use ocamlDoc as soon as possible
- Don't try to program generic functions at the beginning. Only do it when required

Testing

- Write test before programming
- Test only high level behavior
- Plenty of test framework, my preferred one: OUnit
- Do small test or allow to skip big ones (OUnit.skip_if or OUnit.todo)
- Include running test in your default Makefile target

Even if OCaml allow to catch a lot of programming errors, tests is the only way to prevent more subtle error

Optimize

- Never try to optimize an unfinished software
- Extend your tests before beginning to optimize
- Use specific Makefile targets to produce different « flavor »:
 - Bytecode
 - Bytecode with debug
 - Native
 - Native with profiling

**Keep in mind: optimizing any software can be
a long task!**

Optimize

- This is a balance between:
 - Being fast in a particular case
 - Being fast in all cases
 - Spending time to achieve the two points above
- General case: good but very expensive
- Particular case: cheap and cost/time effective
- In fact, this is just as the algorithm complexity:
 - Worst case (general)
 - Average (most used particular case)

Optimize

Optimizing should be about the « average optimization »

- Find your most common cases
- Identify the best algorithm for these cases
- Profile at low level for each case

Optimize (search)

- Find obvious problem:
 - non-terminal recursion
 - inefficient functions called too often
- Find bottleneck
- Find best algorithm
- Find low-level optimization

Optimize

- Optimization is also long work:
 - Time your initial case
 - Do a small change
 - Time your small change
 - Evaluate the win
 - Restart from beginning
- Always try to avoid introducing complicated programming pattern for an optimization
- It is highly improbable that you will be able to optimize a program in a single day

Samples

One very important task when doing optimization is to find your most common cases and transform it into benchmark samples

- Samples should allow you to cover:
 - All your most common cases
 - Different scale of it (from the smaller to the bigger)

Samples

- You should be able to integrate benchmark into your default Makefile target
- Use at least 2 set of benchmark samples:
 - « small »: day to day
 - « big »: nightly/weekly

Samples

- Gathering samples is a complex task:
 - Lack of data
 - Privacy
 - Too small or too big samples
- But OCaml is a fast prototyping language, so just create your sample generator!
 - Adapt to most case
 - Use your customer as seed
 - Include seed and sample generator in your source

Tools

- Don't be mislead by ocamlprof:
 - You must compile everything with the right option to know precisely where you spend time
 - It remains an approximation to what really need to be optimized
 - It has an heavy impact on processing speed, which in turn can change your bottleneck
 - Output is hard to read
- However, ocamlprof is a good tool:
 - For estimating what functions is called often
 - To know what is the time share of the GC

Tools

- A simple trace tools can help you easily
- ocaml-dbug:
 - Stderr output
 - Can be removed at compile time
 - Low impact on performance
- Allow you to find most problem « by eye »:
 - A non-tail recursion implies a quick slowdown in output
 - Time for each part of your program

Moreover, this is also a good tool to debug your program

Tools

- Another good tools is your process data:
 - Memory used
 - Time
 - Temporary file used
- ocaml-process-monitor:
 - Display memory and disk usage
 - Output on stderr (interlaced with ocaml-dbug)

Tools

- Create your benchmark suite:
 - For low level (function call): benchmark OCaml library
 - For high level (whole software on an input): build your own software (using OCaml)
- For low level, integrate your benchmark with your test suite:
 - Only benchmark if test is correct
 - Use the same type of functions (factor your code)
 - Help to dedicate, one test suite to one source module

Tools

For high level, allow to run it through a standard cron job:

- Takes a lot of time to run (15min to several hours)
- Will be launch out of your standard development environment

Benchmarking

- Use a reference implementation and do benchmark relative to it
- Find a way to always produce the same result in a specific context
- Use different scale, different computer architectures and OS
- Keep a log of your benchmark results
- Find a suitable presentation for your results!

Searching for speed

The first step in looking for speed is to understand what is the problem.

- Evaluate difference of performance in your sample and try to know why there is a difference
- Consider asymmetrical evolution of the performance of each sample (why this sample perform better than this one after a change)

Searching for speed (GC)

- Use `ocamlprof` to detect how much the GC is using
- Hook « benchmark » module to monitor GC allocation
- Avoid a lot of fast allocation/deallocation
 - This cause data to migrate from minor to major heap
 - Make the GC works a lot
- Try to reduce GC allocation by reusing datastructure

Searching for speed (GC)

- Avoid maintaining a lot of datastructures for very long time in memory
 - Each time you do a major collection, it will be scanned again and again
- Reuse, when possible, datastructures already allocated (Buffer, String,...)

Searching for speed (imperative)

- Try to use imperative aspect of OCaml when it worth it
- Purely functional programming doesn't apply well to some datastructure (Hashtbl, Array, Queue)
- The limit between purely functional versus imperative is hard to draw:
 - Only incrementing an integer in record is more efficient with purely functional
 - Manipulating string buffer is more efficient with imperative style

Searching for speed

- Non « inlined » function
 - Functor
 - Anonymous function (like in « List.iter »)

Searching for speed (solutions)

- If a lot of datastructure need to be maintained in memory:
 - Use a custom block/C datastructure
 - Use Ancient OCaml module
- A very good imperative datastructure to build your own: Buffer
 - It converges to maximal size of your input data
 - Once at this size, it stops allocating
 - All primitives are good (Buffer.clean, Buffer.reset...)
 - If needed, you can call Buffer.reset to do a « deep »-deallocation

Searching for speed (solutions)

- Play with the different GC call, from minor to compact:
 - Sometimes you don't have a lot of memory allocated, but it is fragmented
 - However, this is only to understand where the problem is, avoid doing this in your final product
- Don't try to play with GC settings, in most case, they are good for general purpose application.

Tips and tricks

- Imperative loop versus recursion
 - No difference of performance “while ...” and “let rec ...”
 - “let rec ..” for complex loop
 - “while ...” to ensure that you have no « non » tail recursion
- The case of int/Int32
 - 31 bits integer is just enough for common case
 - If you need to compute things to 32 bits (or 64 bits), compute using “int” until maximum and switch to Int64/Int32/Big_int when needed

Tips and tricks (32/63 bits)

- OCaml GC rely on pointer
- On amd64:
 - Pointer is twice the size
 - Can address more than 3GB of memory
- Choose what really worth:
 - Memory usage (32)
 - 64 bits number (64)
 - Using more than 3GB memory (64)

Tips and tricks

- Test before change
 - Testing value can be optimized
 - Changing value has always an overhead (`caml_modify`)
- C functions
 - Use C functions for really heavy arithmetic or OS specific task
 - Using C for processing string or unboxed array is almost useless
 - Don't use `caml_register_global_root` if they are many

Tips and tricks

- `ocamlopt` options
 - « `-inline` » has almost no effect if you don't have written a code with this in mind
 - « `-ccopt -O9` » have a real impact on performance
 - « `-unsafe` » and « `-noassert` » only apply to compiled code
- « `unsafe` » functions
 - Whenever you check the bound in a function before processing – use it

OCaml as fast as C!

- OCaml has many advantages:
 - GC is efficient
 - Mixing imperative and functional style can help you to get the best from the two worlds
 - You can make OCaml code run really fast

However, you cannot go as fast as low level hand-made optimization in a language like C.

- But with OCaml:
 - You won't something that works quickly
 - You won't spend time debugging your segfault.

All in all, OCaml is a good bet for speed and cost effective solution

Questions

?